CS 178 Midterm Exam

Machine Learning and Data Mining: Fall 2019

Monday November 4th, 2019

Your name: **Row/Seat Number:**

Your ID #(e.g., 123456789) **UCINetID (e.g.ucinetid@uci.edu)**

- Please put your name and ID **on every page**.
- Total time is 50 minutes. READ THE EXAM FIRST and organize your time; don't spend too long on any one problem.
- Please **write clearly** and **show all your work**.
- If you need clarification on a problem, please raise your hand and wait for the instructor or TA to come over.
- You may use **one** sheet containing handwritten notes for reference, and a (basic) calculator.
- Turn in your notes and any scratch paper with your exam.

Problems

1 Bayes Classifiers, ***(10 points.)*** **3**

2 Nearest Neighbor Regression, ***(12 points.)*** **5**

3 True/False, ***(10 points.)*** **7**

4 Support Vector Machines, ***(10 points.)*** **9**

5 VC-Dimensionality, ***(10 points.)*** **11**

Total, *(52 points.)*

This page is intentionally blank, use as you wish.

Name: ______ ID#: ______

Bayes Classifiers, *(10 points.)*

x_1	x_2	y
c	b	0
b	b	0
b	c	0
a	c	1
a	c	1
a	b	1
a	a	1
b	b	1
c	a	1

Consider the table of measured data given at right. We will use the two observed features x_1 , x_2 to predict the class y. Each feature can take on one of three values, $x_i \in \{a, b, c\}$.
In the case of a tie, we will prefer to predict class $y = 0$.

(1) Write down the probabilities learned by a naïve Bayes classifier: *(4 points.)*

$p(y = 0)$: $p(y = 1)$:

$p(x_1 = a \,|\, y = 0)$: $p(x_1 = a \,|\, y = 1)$:

$p(x_1 = b \,|\, y = 0)$: $p(x_1 = b \,|\, y = 1)$:

$p(x_1 = c \,|\, y = 0)$: $p(x_1 = c \,|\, y = 1)$:

$p(x_2 = a \,|\, y = 0)$: $p(x_2 = a \,|\, y = 1)$:

$p(x_2 = b \,|\, y = 0)$: $p(x_2 = b \,|\, y = 1)$:

$p(x_2 = c \,|\, y = 0)$: $p(x_2 = c \,|\, y = 1)$:

(2) Using your naïve Bayes model, what value of y would you predict given $(x_1 = a, x_2 = b)$?: *(3 points.)*

(3) Using your naïve Bayes model, compute the probabilities: *(3 points.)*
$p(y = 0|x_1 = b, x_2 = c)$: $p(y = 1|x_1 = b, x_2 = c)$:

This page is intentionally blank, use as you wish.

Name: ID#:

Nearest Neighbor Regression, *(12 points.)*

For a regression problem to predict y given a scalar feature x, we observe training data (pictured at right):

x	y
1	3
2.3	3
3.3	1
4	2

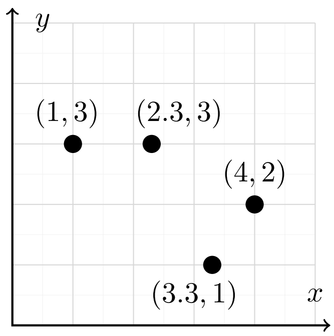

(1) Compute **training** MSE of a 1-nearest neighbor predictor. *(3 points.)*

(2) Compute the **leave-one-out** cross-validation error (MSE) of a 1-nearest neighbor predictor. *(3 points.)*

(3) Compute the **training** MSE of a 2-nearest neighbor predictor. *(3 points.)*

(4) Compute the **leave-one-out** cross-validation error (MSE) of a 2-nearest neighbor predictor. *(3 points.)*

This page is intentionally blank, use as you wish.

Name: ______ ID#: ______

True/False, *(10 points.)*

Here, assume that we have m data points $y^{(i)}$, $x^{(i)}$, $i = 1 \ldots m$, each with n features, $x^{(i)} = [x_1^{(i)} \; \ldots \; x_n^{(i)}]$. For each of the scenarios below, circle one of "true" or "false" to indicate whether you agree with the statement.

True or **false**: In a soft-margin SVM (i.e., loss $\sum_j w_j^2 + R \sum_i \epsilon^{(i)}$), increasing the value of R will make the model more likely to overfit.

True or **false**: A soft-margin SVM model is harder to optimize than a hard-margin SVM, since it is not a quadratic program.

True or **false**: A kernel SVM will be more efficient than a linear SVM when the number of training data, m, is large.

True or **false**: Applying "early stopping" by increasing the convergence tolerance in SGD increases the bias of the learner to reduce overfitting.

True or **false**: When training a perceptron using the logistic negative log-likelihood loss, gradient descent can never become stuck in a local optimum.

True or **false**: Given sufficently many data m, the 1-nearest neighbor classifier error rate approaches the Bayes optimal error rate.

True or **false**: Stochastic gradient descent is often preferred over batch when the number of data points m is very large.

True or **false**: For a perceptron, increasing the regularization penalty of a linear regression model will decrease the resulting model's variance.

True or **false**: For a perceptron, doubling the number of training data available will decrease the resulting model's bias.

True or **false**: For a perceptron, using $2 \times n$ features per data point by adding n random values to each will increase the resulting model's variance.

This page is intentionally blank, use as you wish.

Support Vector Machines, *(10 points.)*

Suppose we are learning a linear support vector machine with two real-valued features x_1, x_2 and binary target $y \in \{-1, +1\}$. We observe training data (pictured at right):

x_1	x_2	y
0.5	2	+1
1	1	+1
1	2	+1
1	3	+1
2	2	-1
3.5	3	-1
4	1	-1

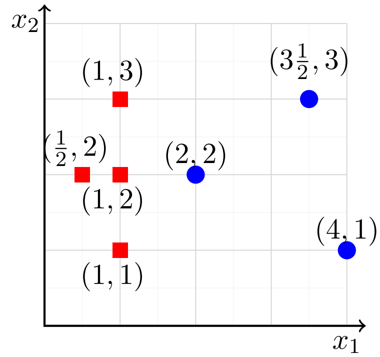

Our linear classifier takes the form

$$f(x; w_1, w_2, b) = \text{sign}(w_1 x_1 + w_2 x_2 + b).$$

(1) Consider the optimal linear SVM classifier for the data, i.e., the one that separates the data and has the largest margin. **Sketch** its decision boundary in the above figure, and **list** the support vectors here. *(2 points.)*

(2) Derive the parameter values w_1, w_2, b of this $f(x)$ using these support vectors. What is the length of the margin? *(3 points.)*

(3) What is the *training error* of a linear SVM on these data? *(2 points.)*

(4) What is the the *leave-one-out cross validation error* for a linear SVM trained on these data? *(3 points.)*

This page is intentionally blank, use as you wish.

Name: ID#:

VC-Dimensionality, *(10 points.)*

Consider the VC dimension of two classifiers defined using two features x_1, x_2.

(1) First, consider a simple classifier f_A that predicts class +1 within a ring with inner radius r and a width of w:

$$f_A(x) = \begin{cases} +1 & (r < (x_1^2 + x_2^2) < r + w) \\ -1 & \text{otherwise} \end{cases}$$

Show that this classifier has VC dimension 2. *(5 points.)*

(2) Now, suppose that we fix $w = 1$, i.e., it is no longer a parameter of the model:

$$f_B(x) = \begin{cases} +1 & (r < (x_1^2 + x_2^2) < r + 1) \\ -1 & \text{otherwise} \end{cases}$$

What is the VC dimension of f_B? Justify your answer.*(5 points.)*

This page is intentionally blank, use as you wish.

This page is intentionally blank, use as you wish.

This page is intentionally blank, use as you wish.

This page is intentionally blank, use as you wish.

This page is intentionally blank, use as you wish.